



---

# ***Chap. 4) Multithreaded Programming***

경희대학교 컴퓨터공학과

방 재 훈

## ■ Heavy-weight

- ✓ A process includes many things:
  - An address space (all the code and data pages)
  - OS resources (e.g., open files) and accounting info.
  - Hardware execution state (PC, SP, registers, etc.)
- ✓ Creating a new process is costly because all of the data structures must be allocated and initialized
  - Linux: over 100 fields in `task_struct` (excluding page tables, etc.)
- ✓ Inter-process communication is costly, since it must usually go through the OS
  - Overhead of system calls and copying data





# Thread Concept: Key Idea

---

- Separate the concept of a process from its execution state
  - ✓ Process: address space, resources, other general process attributes (e.g., privileges)
  - ✓ Execution state: PC, SP, registers, etc.
  - ✓ This execution state is usually called
    - a thread of control,
    - a thread, or
    - a lightweight process (LWP)



# Thread Concept: Key Idea (Cont'd)



# What is a Thread?

---

- A *thread* (or *lightweight process*) is a basic unit of CPU utilization; it consists of:
  - ✓ program counter
  - ✓ register set
  - ✓ stack space
  
- A thread shares with its peer threads its:
  - ✓ code section
  - ✓ data section
  - ✓ operating-system resources
  - ✓ collectively known as a *task* or process
  
- A traditional or *heavyweight* process is equal to a task with one thread

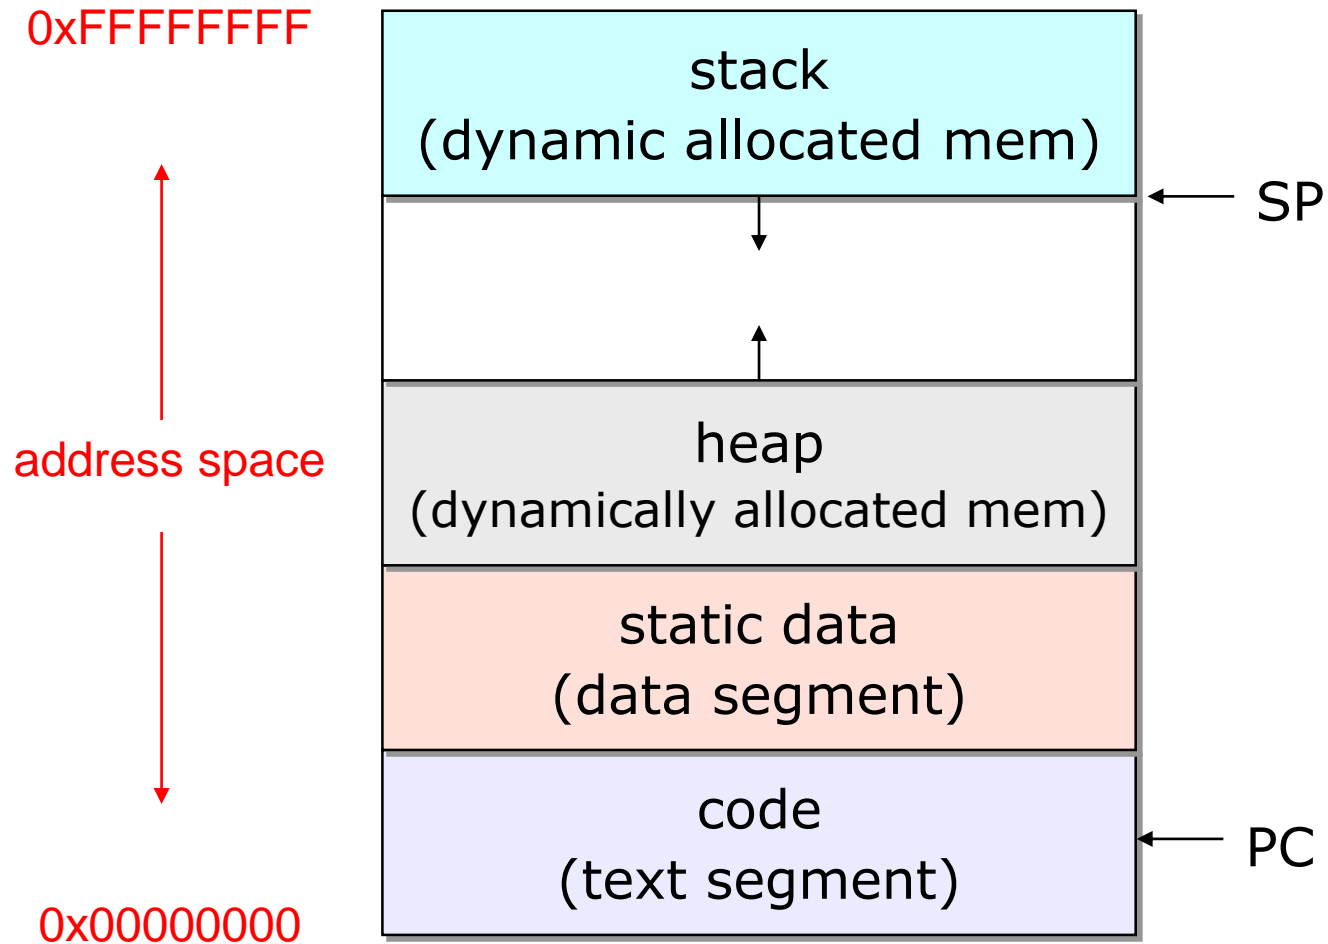


## ■ Processes vs. Threads

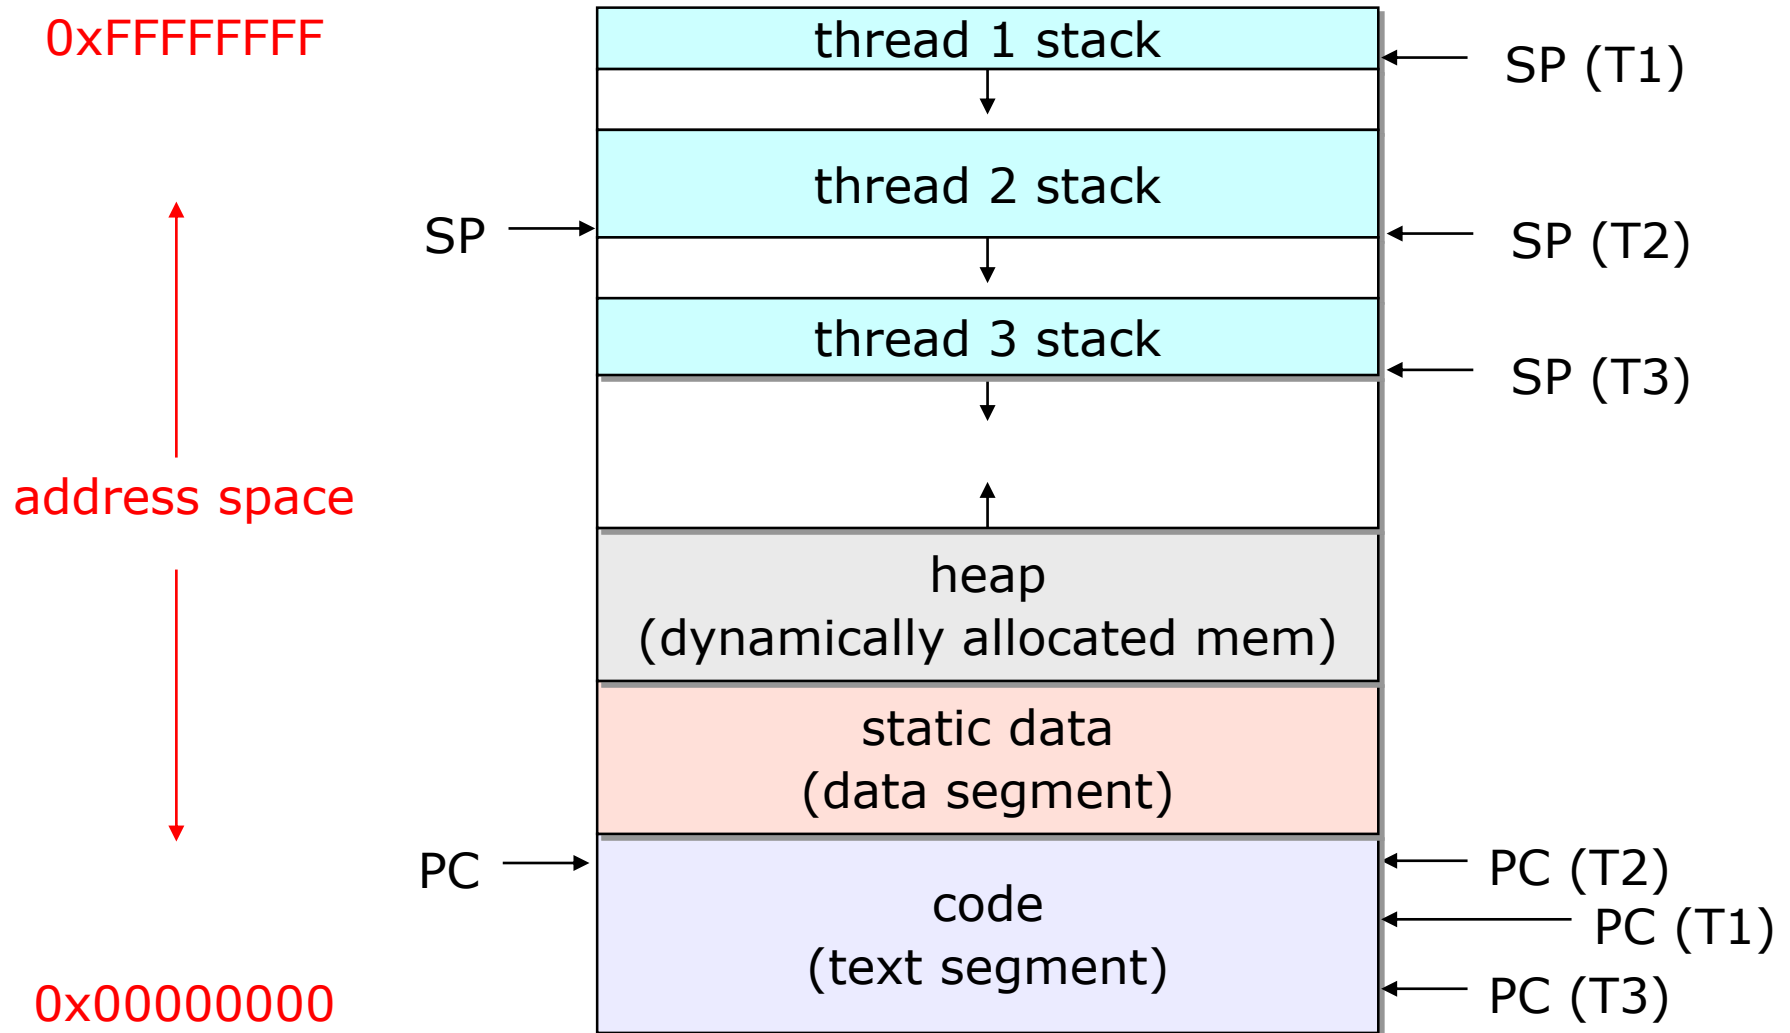
- ✓ A thread is bound to a single process
- ✓ A process, however, can have multiple threads
- ✓ Sharing data between threads is cheap: all see the same address space
- ✓ Threads become the unit of scheduling
- ✓ Processes are now containers in which threads execute
- ✓ Processes become static, threads are the dynamic entities



# Process Address Space

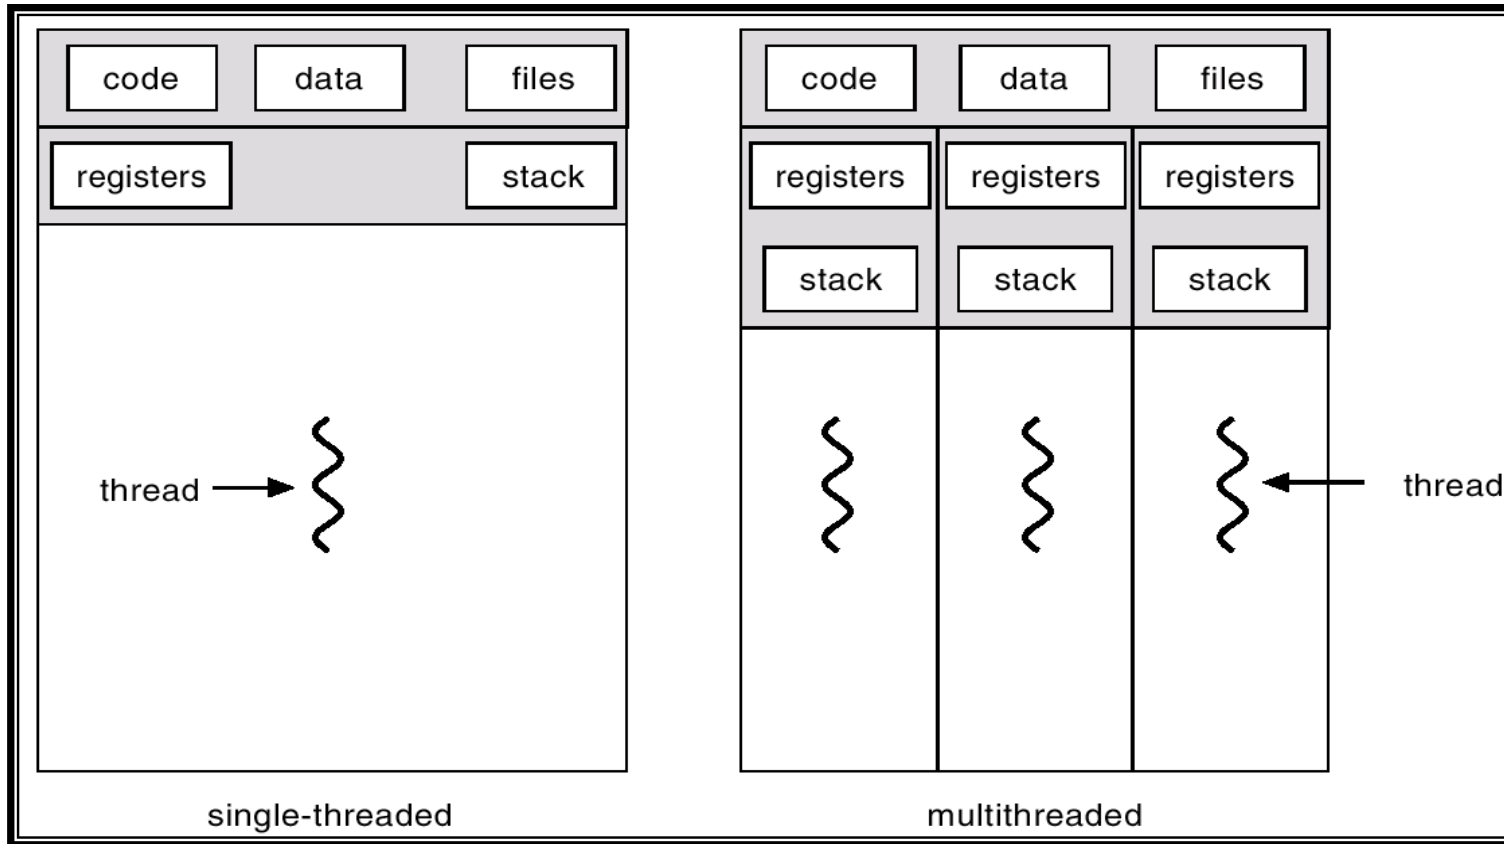


# Address Space with Threads





# Single and Multithreaded Processes



# Concurrent Servers: Processes

---

## ■ Web server example

- ✓ Using fork() to create new processes to handle requests in parallel is overkill for such a simple task.

```
While (1) {  
    int sock = accept();  
    if ((pid = fork()) == 0) {  
        /* Handle client request */  
    } else {  
        /* Close socket */  
    }  
}
```



# Concurrent Servers: Threads

## ■ Using threads

- ✓ We can create a new thread for each request

```
webserver ()
{
    While (1) {
        int sock = accept();
        thread_fork (handle_request, sock);
    }
}
handle_request (int sock)
{
    /* Process request */
    close (sock);
}
```



# ***Benefits***

---

- Responsiveness
- Resource Sharing
- Economy
- Utilization of MP Architectures




# User Threads

- Thread management done by user-level threads library

- Examples

- ✓ POSIX *Pthreads*
- ✓ Mach *C-threads*
- ✓ Solaris *threads*




# Kernel Threads

## ■ Supported by the Kernel

- ✓ thread creation and management requires system calls

## ■ Examples

- ✓ Windows 95/98/NT/2000
- ✓ Solaris
- ✓ Tru64 UNIX
- ✓ BeOS
- ✓ Linux



# User-level Threads vs. Kernel-level Threads

---

## ■ User-level threads

- ✓ The user-level threads library implements thread operations
- ✓ They are small and fast
- ✓ User-level threads are invisible to the OS
- ✓ OS may make poor decisions
  - e.g. blocking I/O
- ✓ Thread scheduling
  - Non-preemptive scheduling: `yield()`
  - Preemptive scheduling: timer through signal

## ■ Kernel-level threads

- ✓ All thread operations are implemented in the kernel
- ✓ The OS schedules all of the threads in a system
- ✓ Kernel threads are cheaper than processes
- ✓ They can still be too expensive



# ***Multithreading Models***

---


- Many-to-One
- One-to-One
- Many-to-Many






# Many-to-One

- Many user-level threads mapped to single kernel thread
- Used on systems that do not support kernel threads




# One-to-One

- Each user-level thread maps to kernel thread
- Examples
  - ✓ Windows 95/98/NT/2000/XP
  - ✓ OS/2



# Many-to-Many Model

- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris 2
- Windows NT/2000/XP with the *ThreadFiber* package



# Threading Issues

---

- Semantics of fork() and exec() system calls
  - ✓ Two versions of fork()
- Thread cancellation
  - ✓ Asynchronous cancellation
  - ✓ Deferred cancellation
- Signal handling
  - ✓ To the thread to which the signal applies
  - ✓ To every thread in the process
  - ✓ To certain threads in the process
  - ✓ Assign a specific thread to receive all signals for the process
- Thread pools
  - ✓ Create a number of threads at process startup
- Thread specific data



# Pthreads

---

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX operating systems



## ■ POSIX-style threads

- ✓ Pthreads
- ✓ DCE threads (early version of Pthreads)
- ✓ Unix International (UI) threads (Solaris threads)
  - Sun Solaris 2, SCO Unixware 2

## ■ Microsoft-style threads

- ✓ Win32 threads
  - Microsoft Windows 98/NT/2000/XP
- ✓ OS/2 threads
  - IBM OS/2



## ■ Thread creation/termination

```
int pthread_create (pthread_t *tid,  
                   pthread_attr_t *attr,  
                   void *(start_routine)(void *),  
                   void *arg);
```

```
void pthread_exit  (void *retval);
```

```
int pthread_join   (pthread_t tid,  
                   void **thread_return);
```



## ■ Mutexes

```
int pthread_mutex_init  
    (pthread_mutex_t *mutex,  
     const pthread_mutexattr_t *mattr);
```

```
int pthread_mutex_destroy  
    (pthread_mutex_t *mutex);
```

```
int pthread_mutex_lock  
    (pthread_mutex_t *mutex);
```

```
int pthread_mutex_unlock  
    (pthread_mutex_t *mutex);
```





## ■ Condition variables

```
int pthread_cond_init  
    (pthread_cond_t *cond,  
     const pthread_condattr_t *cattr);
```

```
int pthread_cond_destroy  
    (pthread_cond_t *cond);
```

```
int pthread_cond_wait  
    (pthread_cond_t *cond,  
     pthread_mutex_t *mutex);
```

```
int pthread_cond_signal  
    (pthread_cond_t *cond);
```

```
int pthread_cond_broadcast  
    (pthread_cond_t *cond);
```



## ■ Thread creation/termination

```
HANDLE CreateThread (lpThreadAttributes, dwStackSize,  
                    lpStartAddress, lpParameter,  
                    dwCreationFlags, lpThreadId);
```

```
void ExitThread (dwExitCode);
```



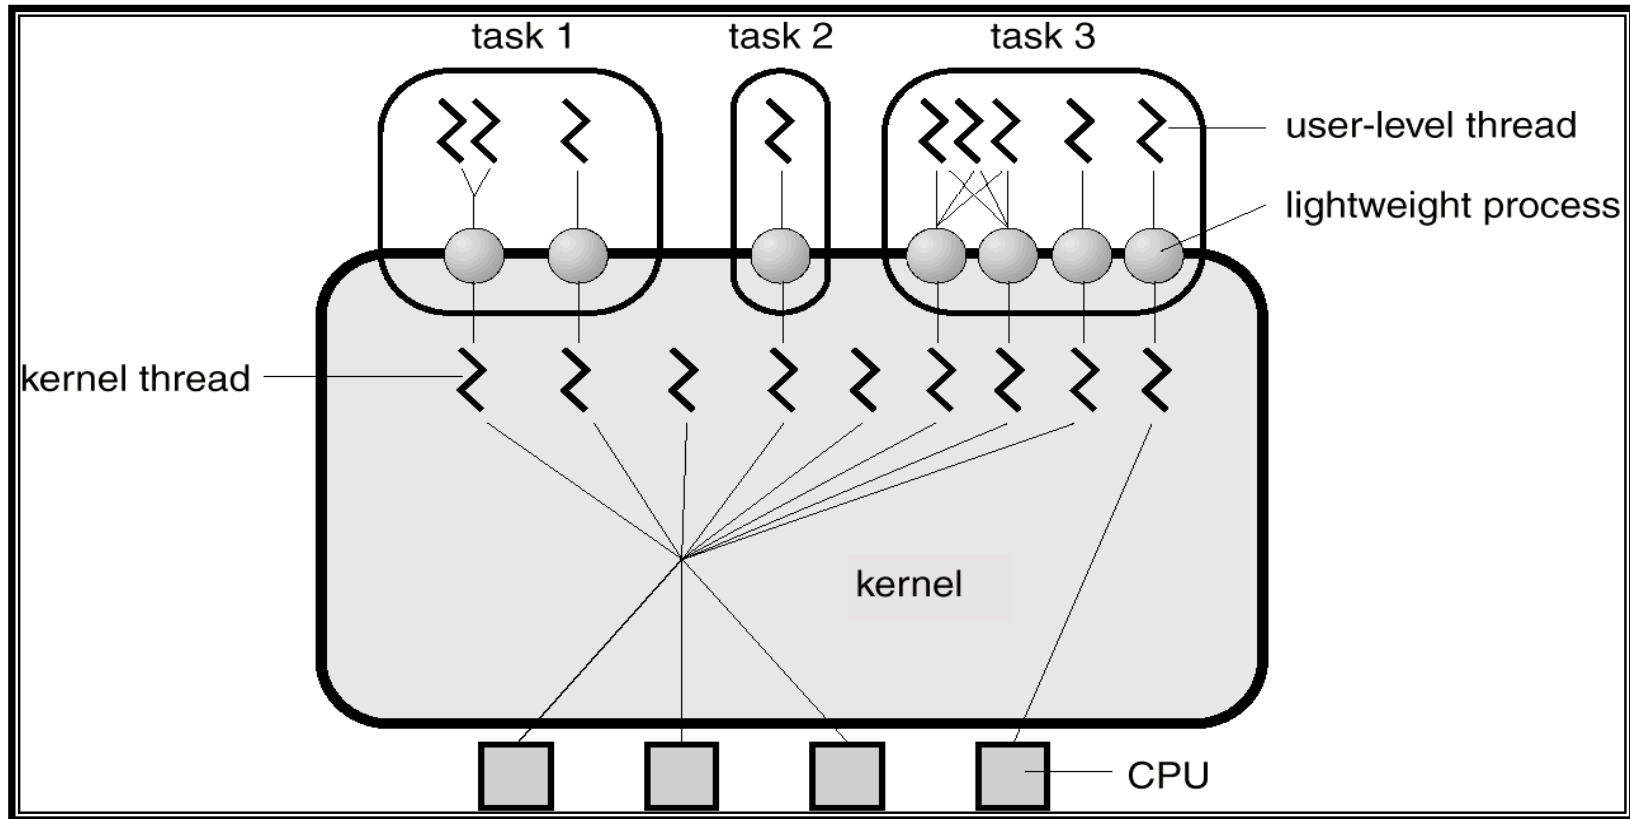
## ■ Thread creation/termination

Create a new class derived from **Thread** class  
Override `run()` method

Create a new class that implements the **Runnable** interface



# Solaris 2 Threads



## ■ LWP (Lightweight Process)

- ✓ A virtual CPU for executing code or system calls
- ✓ Each process contains at least one LWP
- ✓ Each LWP is connected to exactly one kernel-level thread
- ✓ Each LWP is separately dispatched by the kernel, may
  - perform independent system calls
  - incur independent page faults
  - run in parallel on a multiprocessor, etc.
- ✓ The thread library dynamically adjusts the number of LWPs in the pool to ensure the best performance for the application
- ✓ It also “ages” LWPs and deletes them when they are unused for a long time.
- ✓ An LWP is a kernel data structure

## ■ *For implementing many-to-many model*



# Windows XP Threads

---

- Implements the one-to-one mapping
- Each thread contains
  - ✓ a thread id
  - ✓ register set
  - ✓ separate user and kernel stacks
  - ✓ private data storage area
- Cf) Fibers
  - ✓ Fibers are often called “lightweight” threads
  - ✓ Fibers are invisible to the kernel
  - ✓ Fibers provide a functionality of the many-to-many model

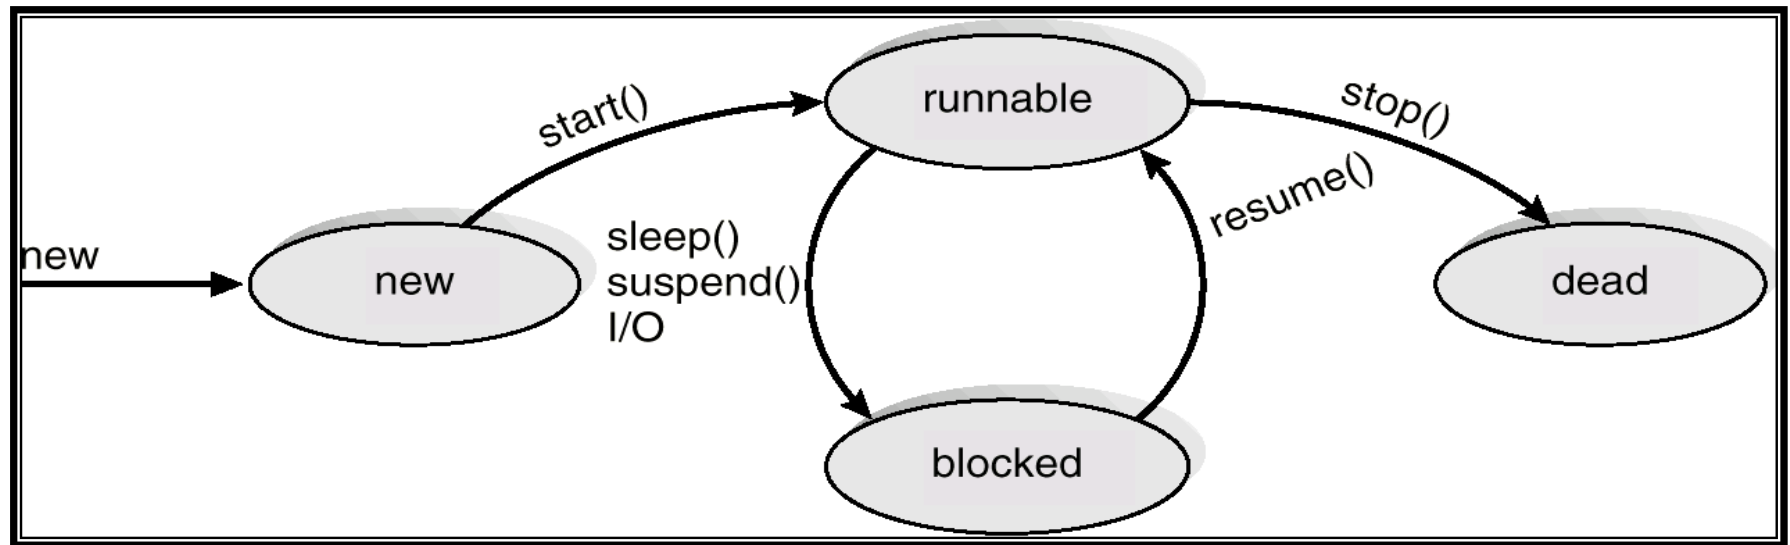


- Linux 2.4 introduces a concept of “thread groups”
  - ✓ *tasks* rather than *threads*
  - ✓ Thread creation is done through clone() system call
  - ✓ Clone() allows a child task to share the address space of the parent task (process)
  - ✓ So, there exist POSIX compatibility problems
  
- Approaches for POSIX compliance
  - ✓ NPTL (Native POSIX Threading Library) – by RedHat
    - 1:1 model
  - ✓ NGPT (Next Generation POSIX Threading) – by IBM
    - M:N model → Linux 2.6



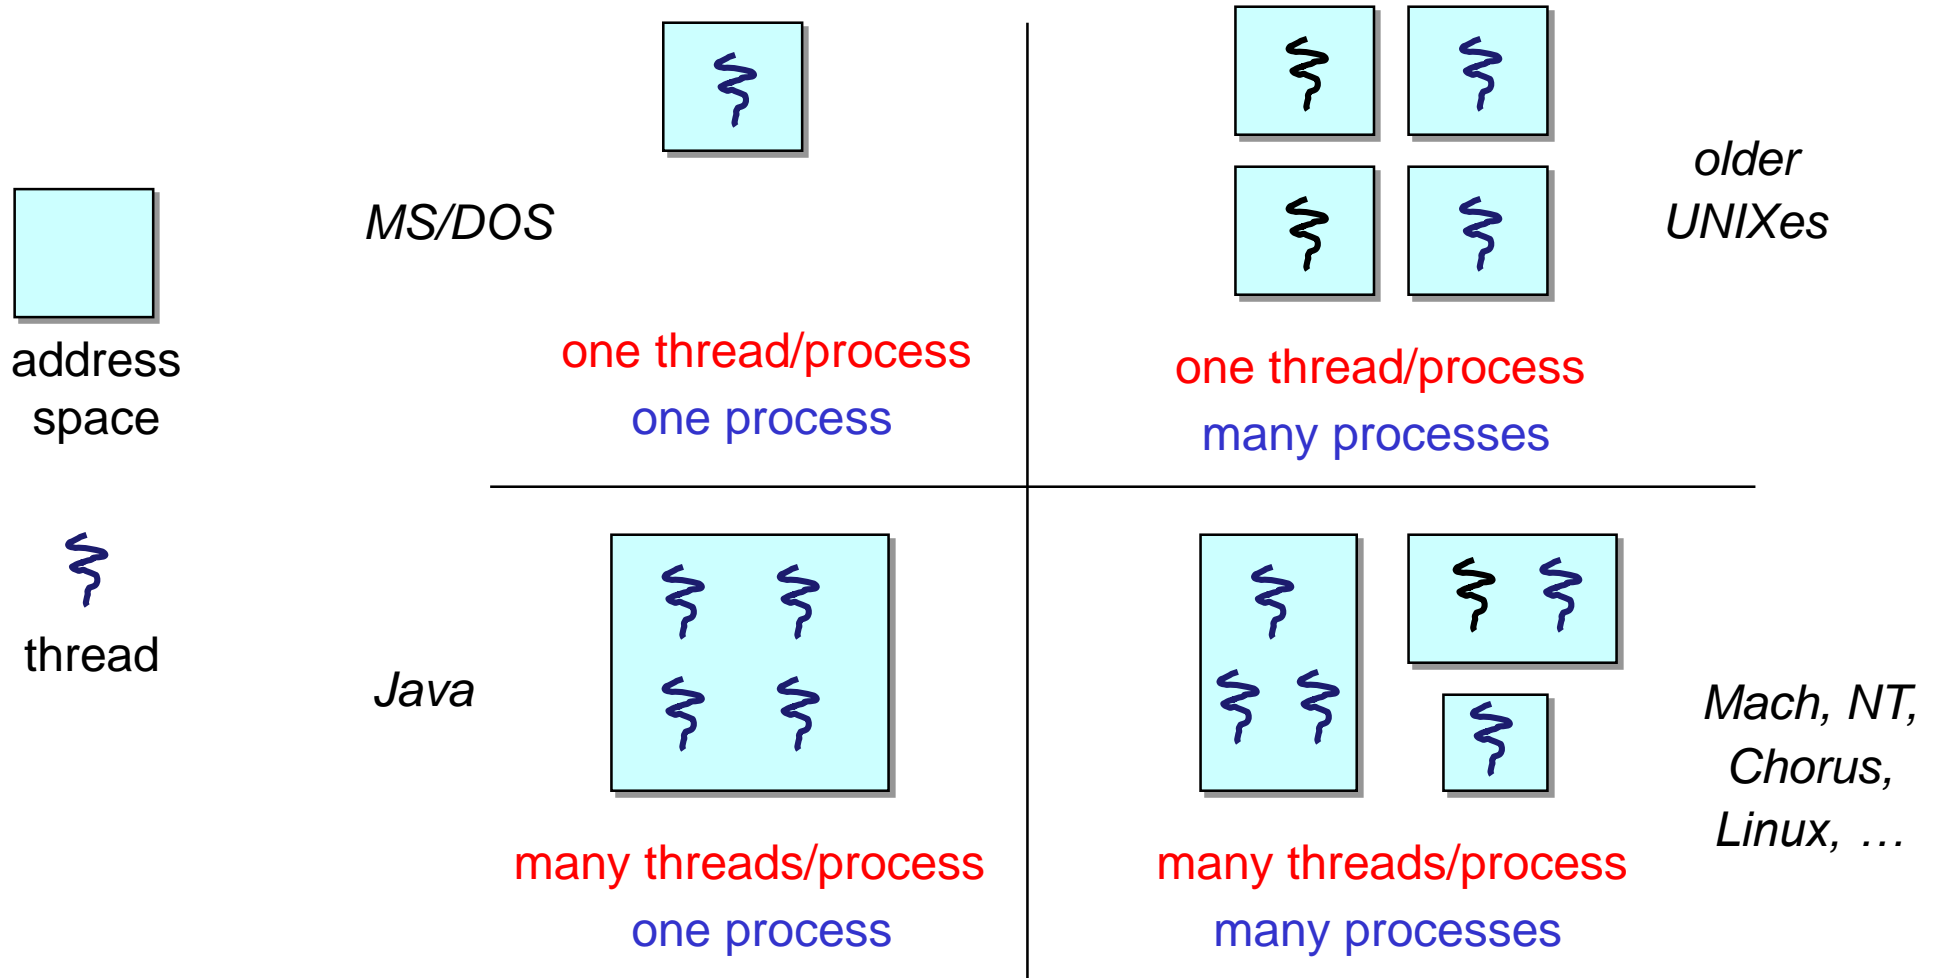
# Java Threads

- Java threads may be created by:
  - ✓ Extending Thread class
  - ✓ Implementing the Runnable interface
- Java threads are managed by the JVM
- Java thread states





# Threads Design Space



## ■ Thread concept

- ✓ Separate the concept of a process from its execution state
- ✓ Execution state: PC, SP, registers, etc.

## ■ Multithreading models

- ✓ User threads to kernel threads mapping
- ✓ Many-to-one
  - User-level threads implementation
- ✓ One-to-one
  - MS-Windows
- ✓ Many-to-many
  - Solaris, Linux

## ■ Multithreaded programming

- ✓ Unix, Linux: POSIX-style threads (Pthread API)
- ✓ MS-Windows: Win32 threads (Win32 API)
- ✓ Java: Java threads (Thread class)

