



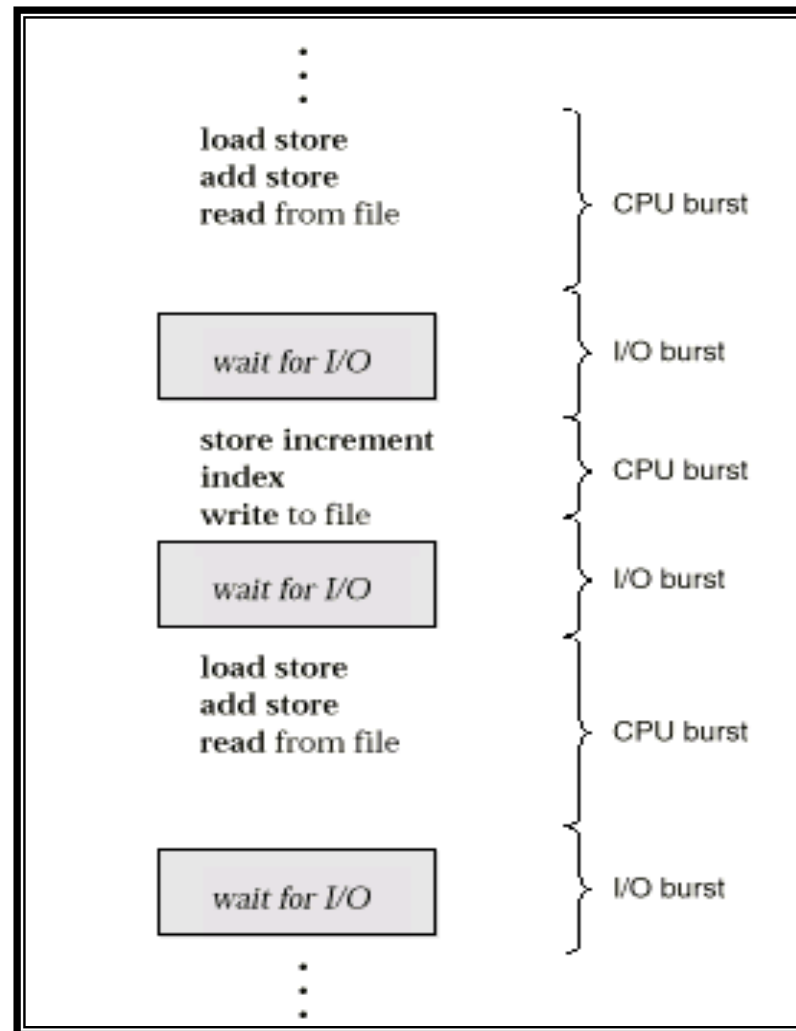
---

# ***Chap. 5) Process Scheduling***

경희대학교 컴퓨터공학과

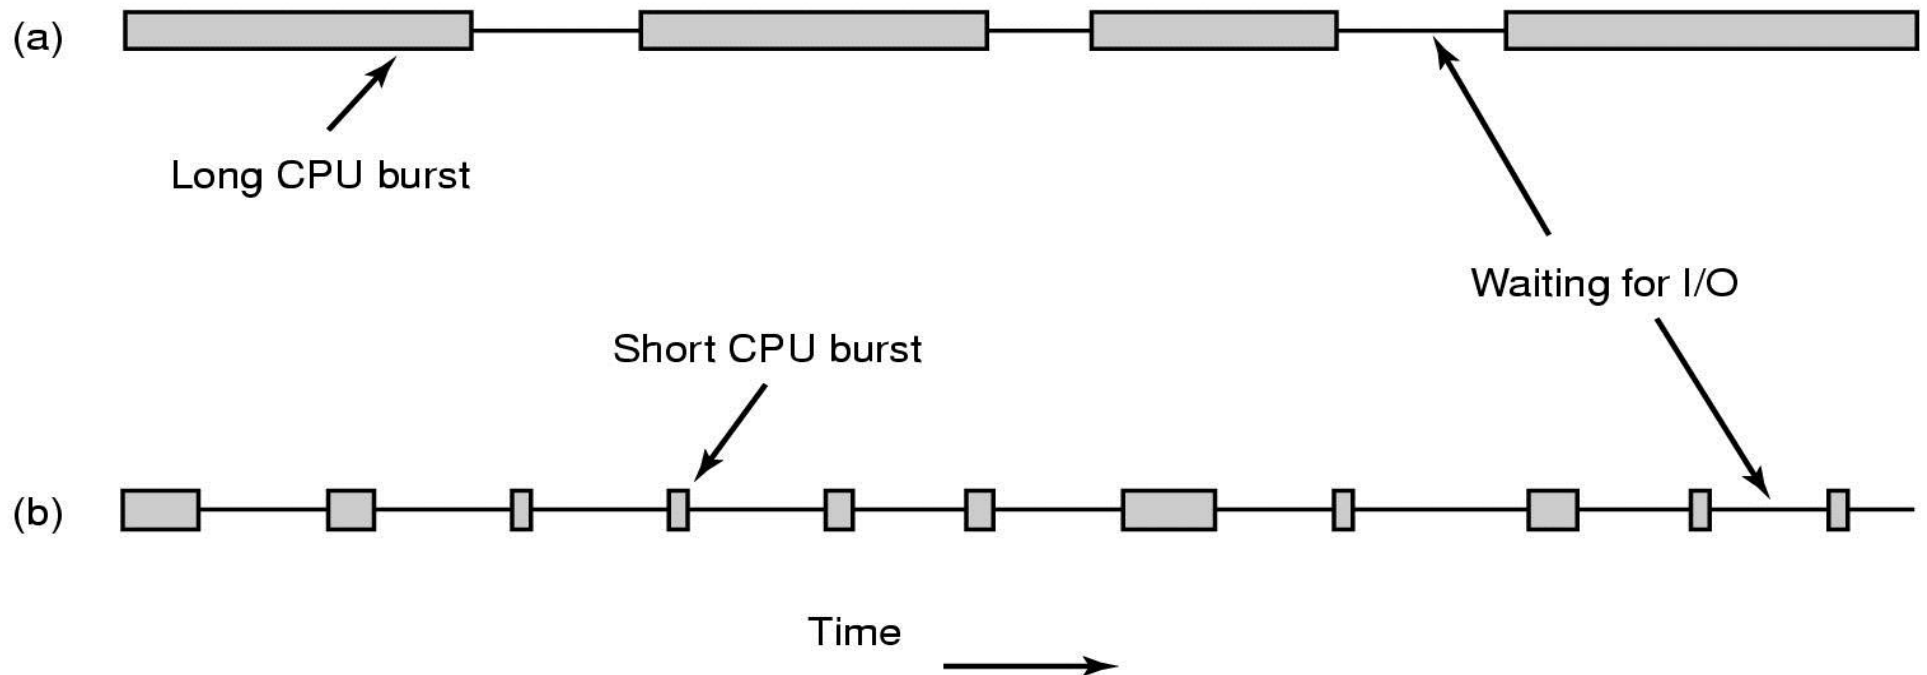
방 재 훈

# Alternating Sequence of CPU And I/O Bursts



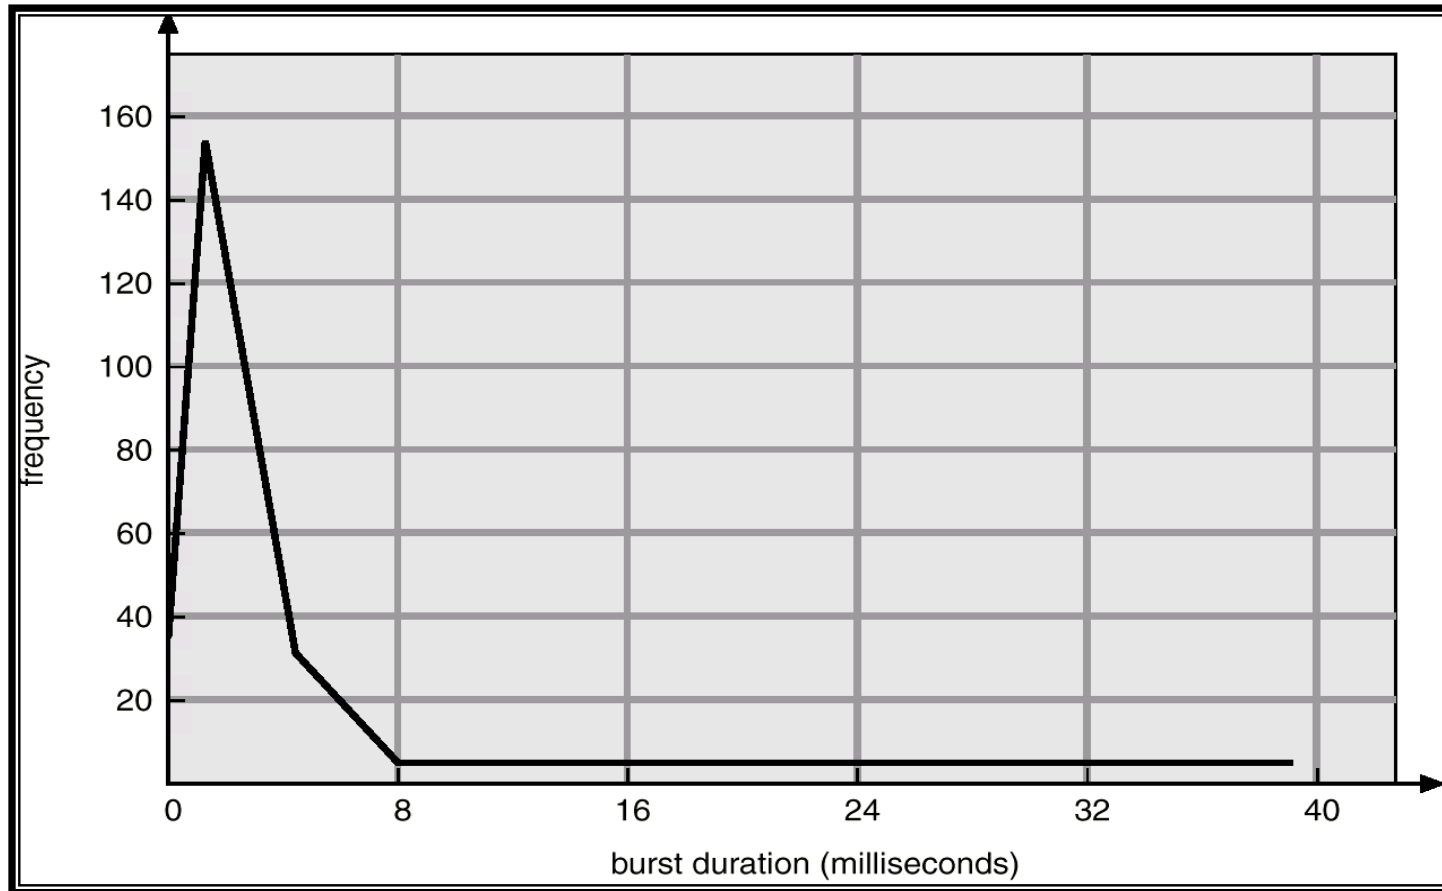
# CPU burst vs. I/O burst

- (a) A CPU-bound process
- (b) An I/O-bound process



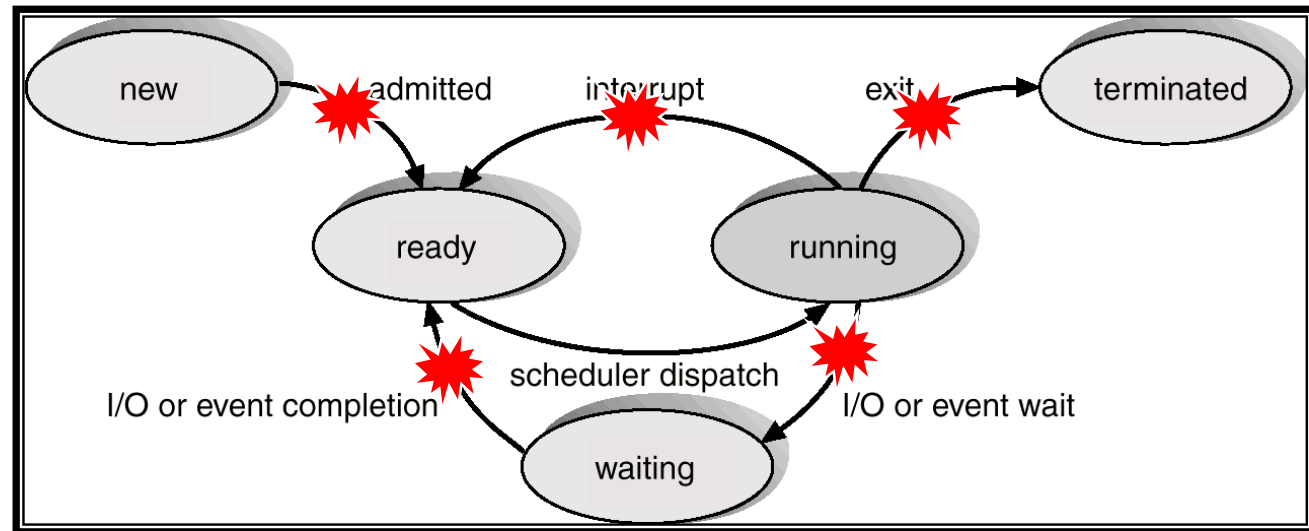
# Histogram of CPU-burst Times

---



# CPU Scheduler

- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them
- CPU scheduling decisions may take place when a process:
  1. Switches from running to waiting state
  2. Switches from running to ready state
  3. Switches from waiting to ready
  4. Terminates



- Scheduling under 1 and 4 is *nonpreemptive*
- All other scheduling is *preemptive*



# Dispatcher

---

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
  - ✓ switching context
  - ✓ switching to user mode
  - ✓ jumping to the proper location in the user program to restart that program
  
- *Dispatch latency*
  - ✓ time it takes for the dispatcher to stop one process and start another running



# Preemptive vs. Non-preemptive

---

## ■ Non-preemptive scheduling

- ✓ The scheduler waits for the running job to explicitly (voluntarily) block
- ✓ Scheduling takes place only when
  - A process switches from running to waiting state
  - A process terminates

## ■ Preemptive scheduling

- ✓ The scheduler can interrupt a job and force a context switch
- ✓ What happens
  - If a process is preempted in the midst of updating the shared data?
  - If a process in system call is preempted?



# Scheduling Criteria

---

## ■ CPU utilization

- ✓ keep the CPU as busy as possible

## ■ Throughput

- ✓ # of processes that complete their execution per time unit

## ■ Turnaround time

- ✓ amount of time to execute a particular process

## ■ Waiting time

- ✓ amount of time a process has been waiting in the ready queue

## ■ Response time

- ✓ amount of time it takes from when a request was submitted until the first response is produced, **not** output (for time-sharing environment)





# *Optimization Criteria*

---

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time



## ■ All systems

- ✓ **Fairness**: giving each process a fair share of the CPU
- ✓ **Balance**: keeping all parts of the system busy

## ■ Batch systems

- ✓ **Throughput**: maximize jobs per hour
- ✓ **Turnaround time**: minimize time between submission and termination
- ✓ **CPU utilization**: keep the CPU busy all the time



# Scheduling Goals (Cont'd)

---

## ■ Interactive systems

- ✓ **Response time**: minimize average time spent on wait queue
- ✓ **Waiting time**: minimize average time spent on ready queue
- ✓ **Proportionality**: meet users' expectations

## ■ Real-time systems

- ✓ **Meeting deadlines**: avoid losing data
- ✓ **Predictability**: avoid quality degradation in multimedia systems



## ■ Starvation

- ✓ A situation where a process is prevented from making progress because another process has the resource it requires.
  - Resource could be the CPU or a lock
- ✓ A poor scheduling policy can cause starvation
  - If a high-priority process always prevents a low-priority process from running on the CPU
- ✓ Synchronization can also cause starvation
  - One thread always beats another when acquiring a lock
  - Constant supply of readers always blocks out writers



## ■ First-Come, First-Served

- ✓ Jobs are scheduled in order that they arrive
- ✓ “Real-world” scheduling of people in lines
  - e.g. supermarket, bank tellers, McDonalds, etc.
- ✓ Typically, non-preemptive
- ✓ Jobs are treated equally: no starvation

## ■ Problems


- ✓ Average waiting time can be large if small jobs wait behind long ones
  - Basket vs. cart
- ✓ May lead to poor overlap of I/O and CPU

# First-Come, First-Served (FCFS) Scheduling

---

<u>Process</u>	<u>Burst Time</u>
$P_1$	24
$P_2$	3
$P_3$	3

- Suppose that the processes arrive in the order:  $P_1$ ,  $P_2$ ,  $P_3$   
The Gantt Chart for the schedule is:



- Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$
- Average waiting time:  $(0 + 24 + 27)/3 = 17$



# FCFS Scheduling (Cont'd)

---

- Suppose that the processes arrive in the order

$$P_2, P_3, P_1.$$

- The Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$
- Average waiting time:  $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- *Convoy effect*
  - ✓ short process behind long process



## ■ Shortest Job First

- ✓ Choose the job with the smallest expected CPU burst
- ✓ Can prove that SJF has optimal min. average waiting time
  - Only when all jobs are available simultaneously
- ✓ Non-preemptive

## ■ Problems

- ✓ Impossible to know size of future CPU burst
- ✓ Can you make a reasonable guess?
- ✓ Can potentially starve



# Shortest-Job-First (SJF) Scheduling

---


- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
  - ✓ Nonpreemptive
    - Once CPU given to the process it cannot be preempted until completes its CPU burst
  - ✓ Preemptive
    - If a new process arrives with CPU burst length less than remaining time of current executing process, preempt
    - This scheme is known as the Shortest-Remaining-Time-First (SRTF)
- SJF is *optimal*
  - ✓ gives minimum average waiting time for a given set of processes



# Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

## ■ SJF (non-preemptive)




## ■ Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$



# Example of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

## ■ SJF (preemptive) (= SRTF)



## ■ Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$



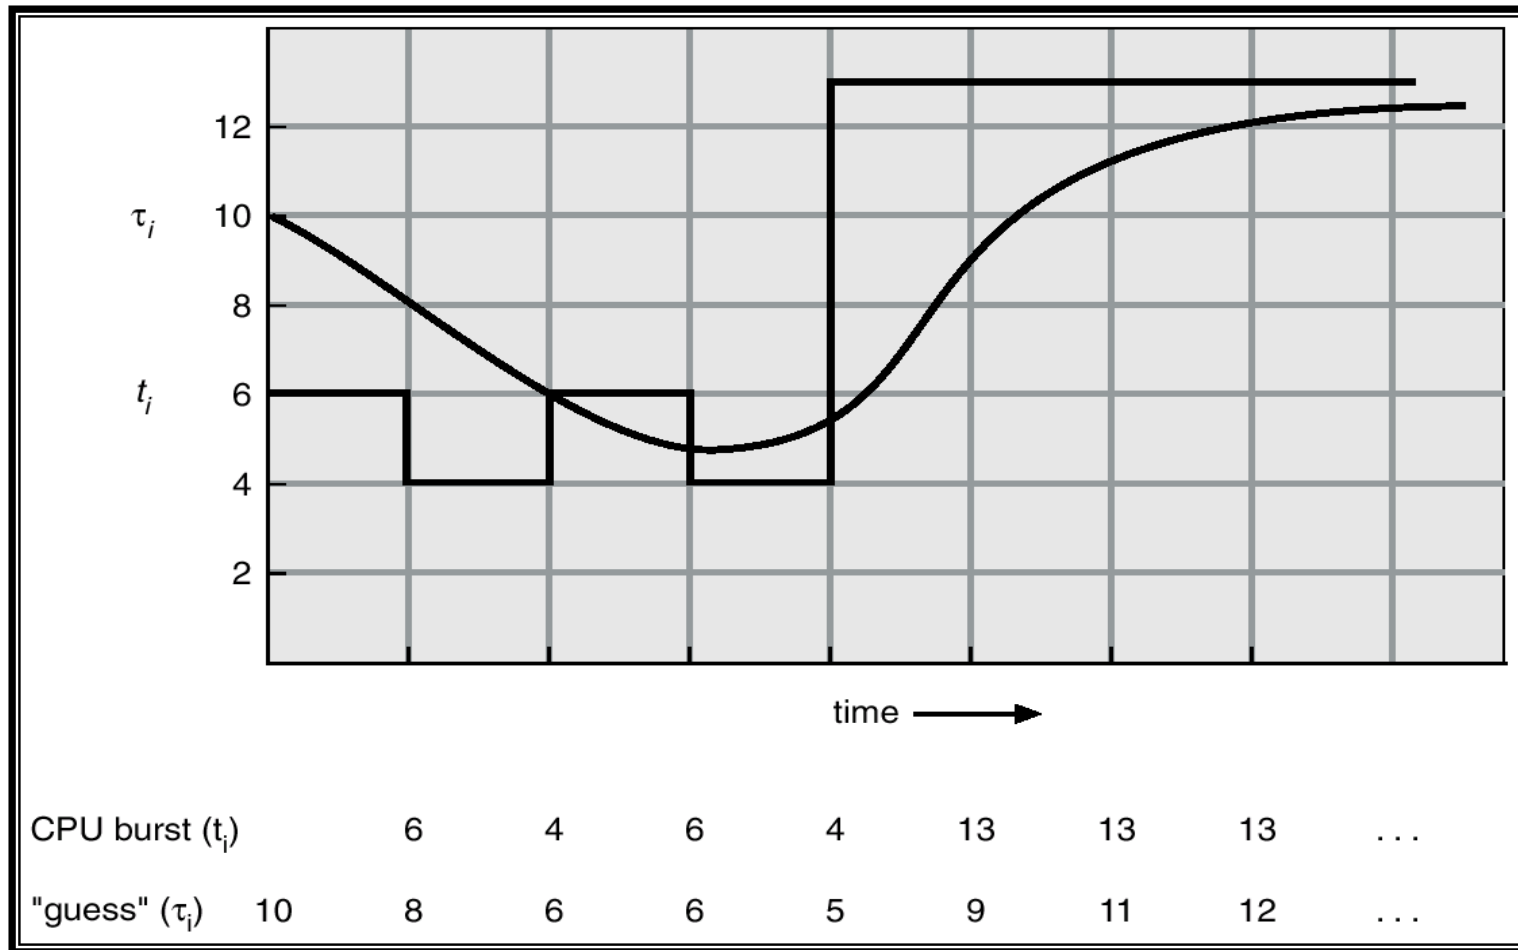
# Determining Length of Next CPU Burst

---

- Can only estimate the length
- Can be done by using the length of previous CPU bursts, using exponential averaging
  1.  $t_n$  = actual length of  $n^{th}$  CPU burst
  2.  $\tau_{n+1}$  = predicted value for the next CPU burst
  3.  $\alpha, 0 \leq \alpha \leq 1$
  4. Define:  $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$ .



# Prediction of the Length of the Next CPU Burst



# Examples of Exponential Averaging

---

## ■ $\alpha = 0$

- ✓  $\tau_{n+1} = \tau_n$
- ✓ Recent history does not count

## ■ $\alpha = 1$

- ✓  $\tau_{n+1} = t_n$
- ✓ Only the actual last CPU burst counts

## ■ If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n-1} t_0\end{aligned}$$

- Since both  $\alpha$  and  $(1 - \alpha)$  are less than or equal to 1, each successive term has less weight than its predecessor



# Priority Scheduling

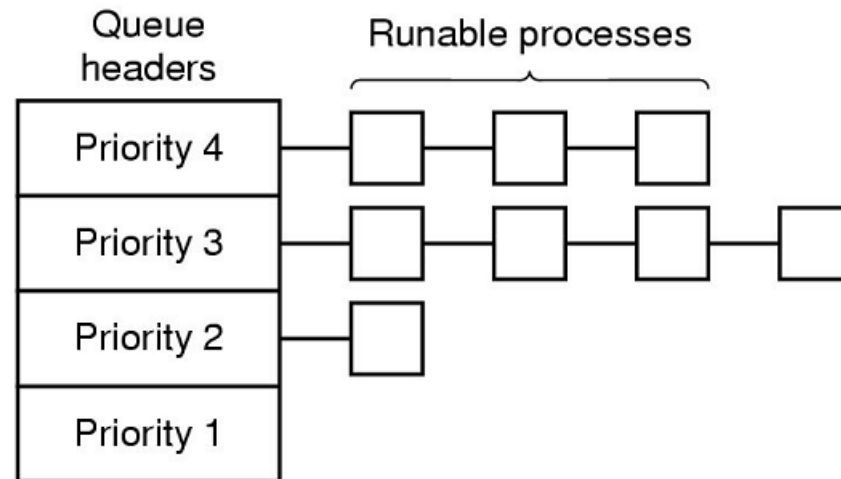
---

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority)
  - ✓ Preemptive
  - ✓ Nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- Problem  $\equiv$  Starvation (or Indefinite blocking)
  - ✓ low priority processes may never execute
- Solution  $\equiv$  Aging
  - ✓ as time progresses increase the priority of the process



# Priority Scheduling

- Abstractly modeled as multiple “priority queues”
  - ✓ Put ready job on Q associated with its priority





# Round Robin (RR)

---

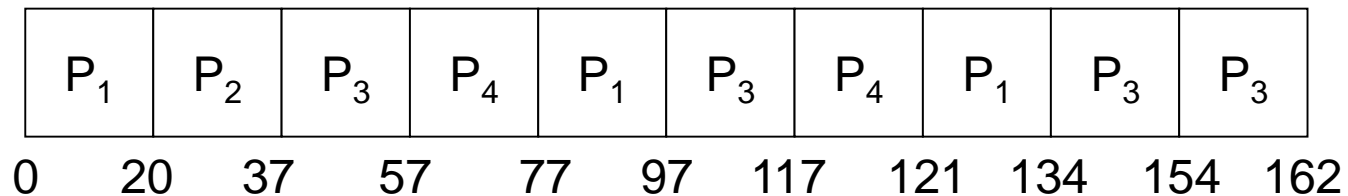
- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds
  - ✓ After this time has elapsed, the process is preempted and added to the end of the ready queue
  
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once
  - ✓ No process waits more than  $(n-1)q$  time units
  
- Performance
  - ✓  $q$  large  $\Rightarrow$  FIFO
  - ✓  $q$  small  $\Rightarrow$   $q$  must be large with respect to context switch, otherwise overhead is too high



## Example of RR with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
$P_1$	53
$P_2$	17
$P_3$	68
$P_4$	24

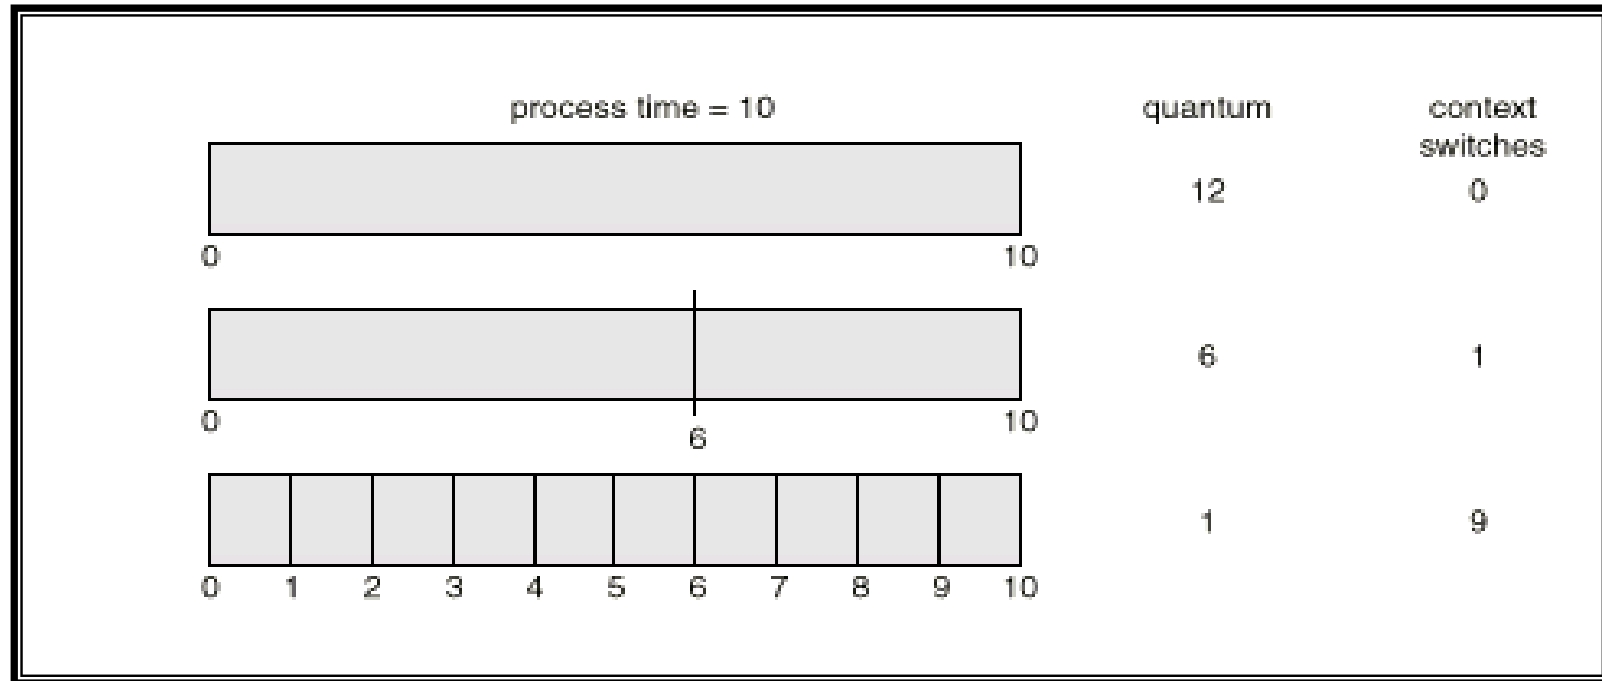
- The Gantt chart is:



- Typically, higher average turnaround than SJF, but better *response*



# Time Quantum and Context Switch Time



## ■ What do you set the quantum to be?

- ✓ quantum  $\rightarrow \infty$  : FIFO
- quantum  $\rightarrow 0$  : processor sharing
- ✓ If small, then context switches are frequent incurring high overhead (CPU utilization drops)
- ✓ If large, then response time drops
- ✓ A rule of thumb: 80% of the CPU bursts should be shorter than the time quantum

## ■ Treats all jobs equally

- ✓ Multiple background jobs?



# Combining Algorithms

---

- Scheduling algorithms can be combined in practice
  - ✓ Have multiple queues
  - ✓ Pick a different algorithm for each queue
  - ✓ Have a mechanism to schedule among queues
  - ✓ And maybe, move processes between queues



# Multilevel Queue

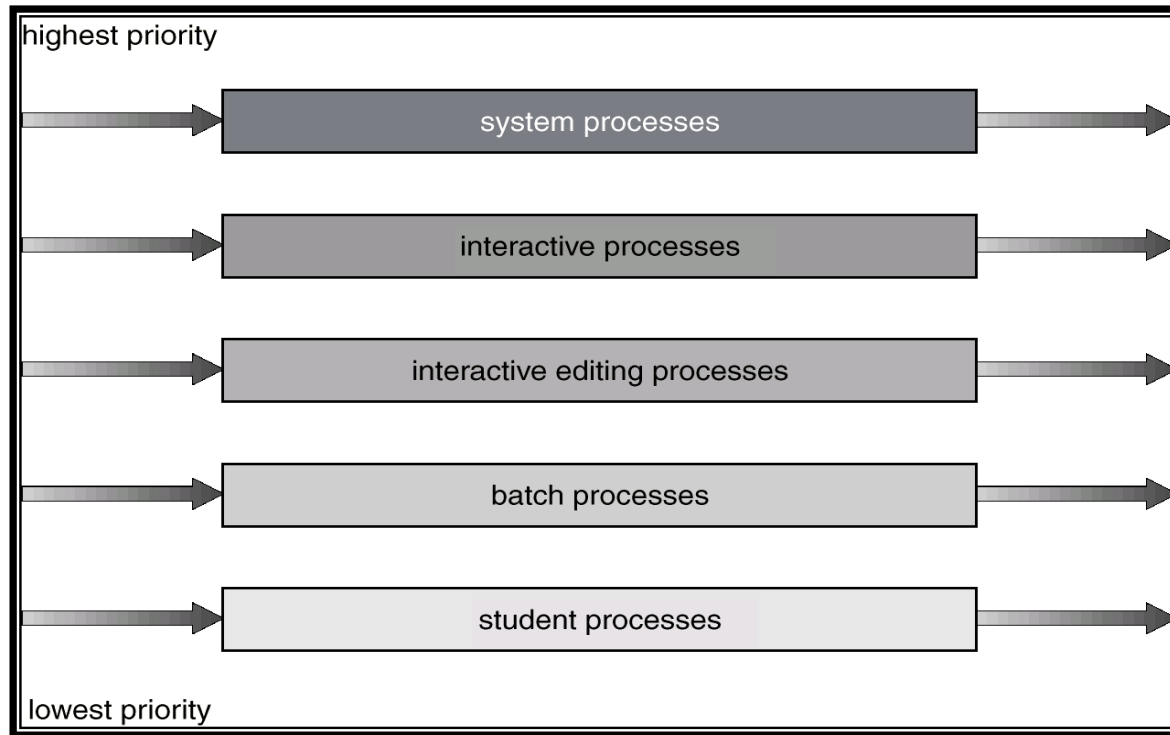
---

- Ready queue is partitioned into separate queues:
  - ✓ foreground (interactive)
  - ✓ background (batch)
  
- Each queue has its own scheduling algorithm:
  - ✓ foreground – RR
  - ✓ background – FCFS
  
- Scheduling must be done between the queues
  - ✓ Fixed priority scheduling
    - (i.e., serve all from foreground then from background) Possibility of starvation
  - ✓ Time slice
    - each queue gets a certain amount of CPU time which it can schedule amongst its processes
    - i.e., 80% to foreground in RR & 20% to background in FCFS



# Multilevel Queue Scheduling

---



# Multilevel Feedback Queue

---

- A process can move between the various queues
  - ✓ aging can be implemented this way
  
- Multilevel-feedback-queue scheduler defined by the following parameters:
  - ✓ number of queues
  - ✓ scheduling algorithms for each queue
  - ✓ method used to determine when to upgrade a process
  - ✓ method used to determine when to demote a process
  - ✓ method used to determine which queue a process will enter when that process needs service





# Example of Multilevel Feedback Queue

---

## ■ Three queues:

- ✓  $Q_0$  – time quantum 8 milliseconds
- ✓  $Q_1$  – time quantum 16 milliseconds
- ✓  $Q_2$  – FCFS

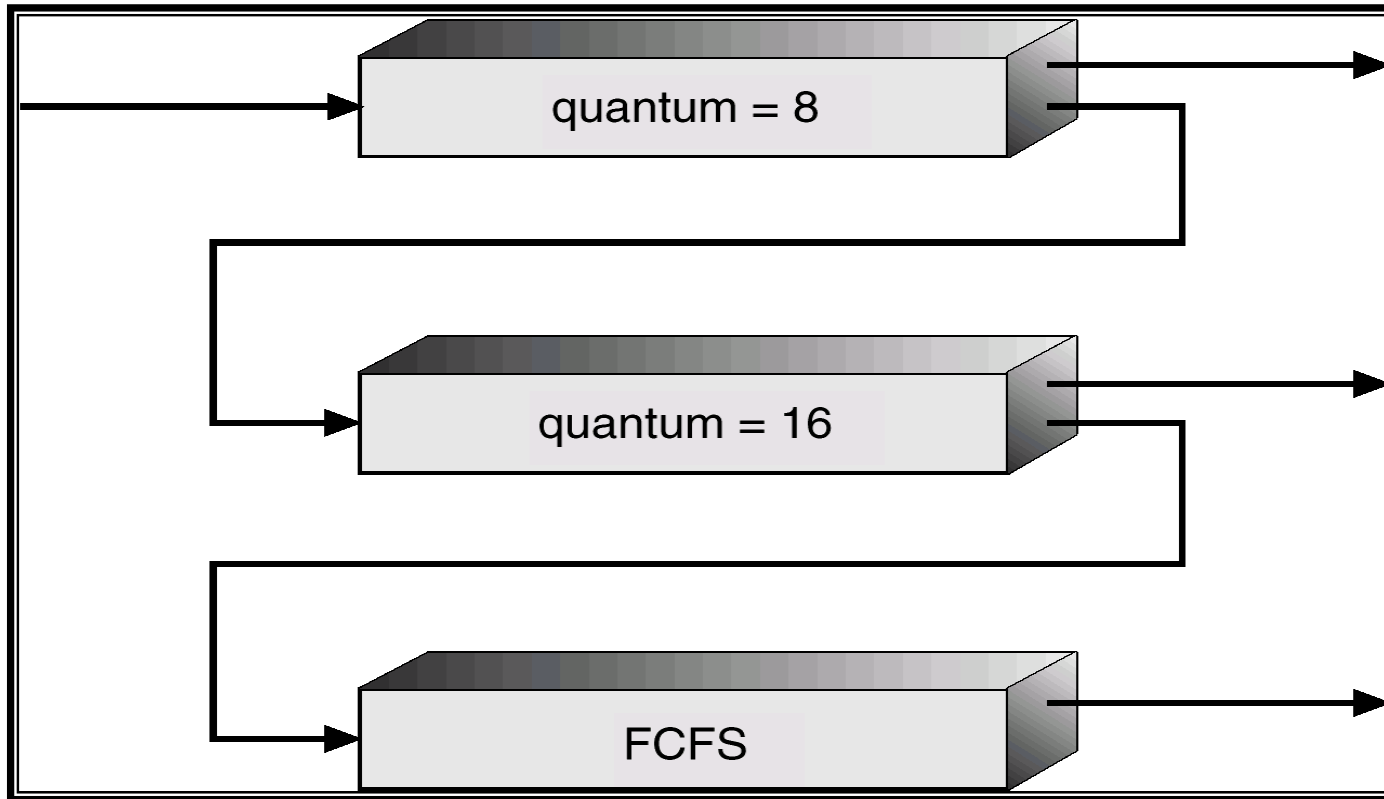
## ■ Scheduling

- ✓ A new job enters queue  $Q_0$  which is served FCFS
- ✓ When it gains CPU, job receives 8 milliseconds
- ✓ If it does not finish in 8 milliseconds, job is moved to queue  $Q_1$
- ✓ At  $Q_1$  job is again served FCFS and receives 16 additional milliseconds
- ✓ If it still does not complete, it is preempted and moved to queue  $Q_2$



# Multilevel Feedback Queues

---



## ■ The canonical UNIX scheduler uses a MLFQ

- ✓ 3 – 4 classes spanning ~170 priority levels
  - Timeshare, System, Real-time, Interrupt (Solaris 2)
- ✓ Priority scheduling across queues, RR within a queue
  - The process with the highest priority always runs
  - Processes with the same priority are scheduled RR
- ✓ Processes dynamically change priority
  - Increases over time if process blocks before end of quantum
  - Decreases over time if process uses entire quantum

## ■ Motivation

- ✓ The idea behind the UNIX scheduler is to reward interactive processes over CPU hogs
- ✓ Interactive processes typically run using short CPU bursts
  - They do not finish quantum before waiting for more input
- ✓ Want to minimize response time
  - Time from keystroke (putting process on ready queue) to executing the handler (process running)
  - Don't want editor to wait until CPU hog finishes quantum
- ✓ This policy delays execution of CPU-bound jobs



# Multiple-Processor Scheduling

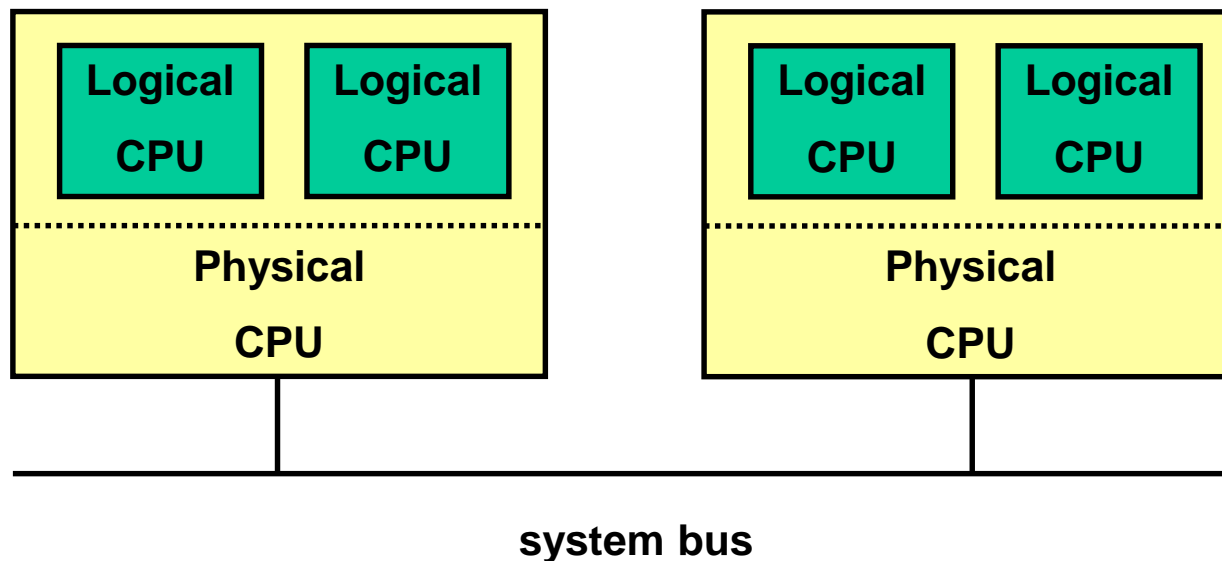
---

- CPU scheduling more complex when multiple CPUs are available
- *Homogeneous processors within a multiprocessor*
  - ✓ UMA (Uniform Memory Access)
- *Load sharing*
- *Asymmetric multiprocessing*
  - ✓ Only one processor accesses the system data structures, alleviating the need for data sharing
  - ✓ Not efficient



# Multiple-Processor Scheduling

- Symmetric multithreading (SMT)
- Hyperthreading technology on Intel Processors
- Typical SMT architecture



# Real-Time Scheduling

---

## ■ *Hard real-time systems*

- ✓ required to complete a critical task within a guaranteed amount of time

## ■ *Soft real-time systems*

- ✓ requires that critical processes receive priority over less fortunate ones

## ■ *Static vs. Dynamic priority scheduling*

- ✓ Static: Rate-Monotonic algorithm
- ✓ Dynamic: EDF (Earliest Deadline First) algorithm



# Algorithm Evaluation

---

## ■ Deterministic modeling

- ✓ Takes a particular predetermined workload and defines the performance of each algorithm for that workload

## ■ Queueing models

- ✓ Mathematical models used to compute expected system parameters

## ■ Simulation

- ✓ Algorithmic models which simulate a simplified version of a system using statistical input
- ✓ Trace tape (or trace data)
- ✓ *Cf) Emulation*

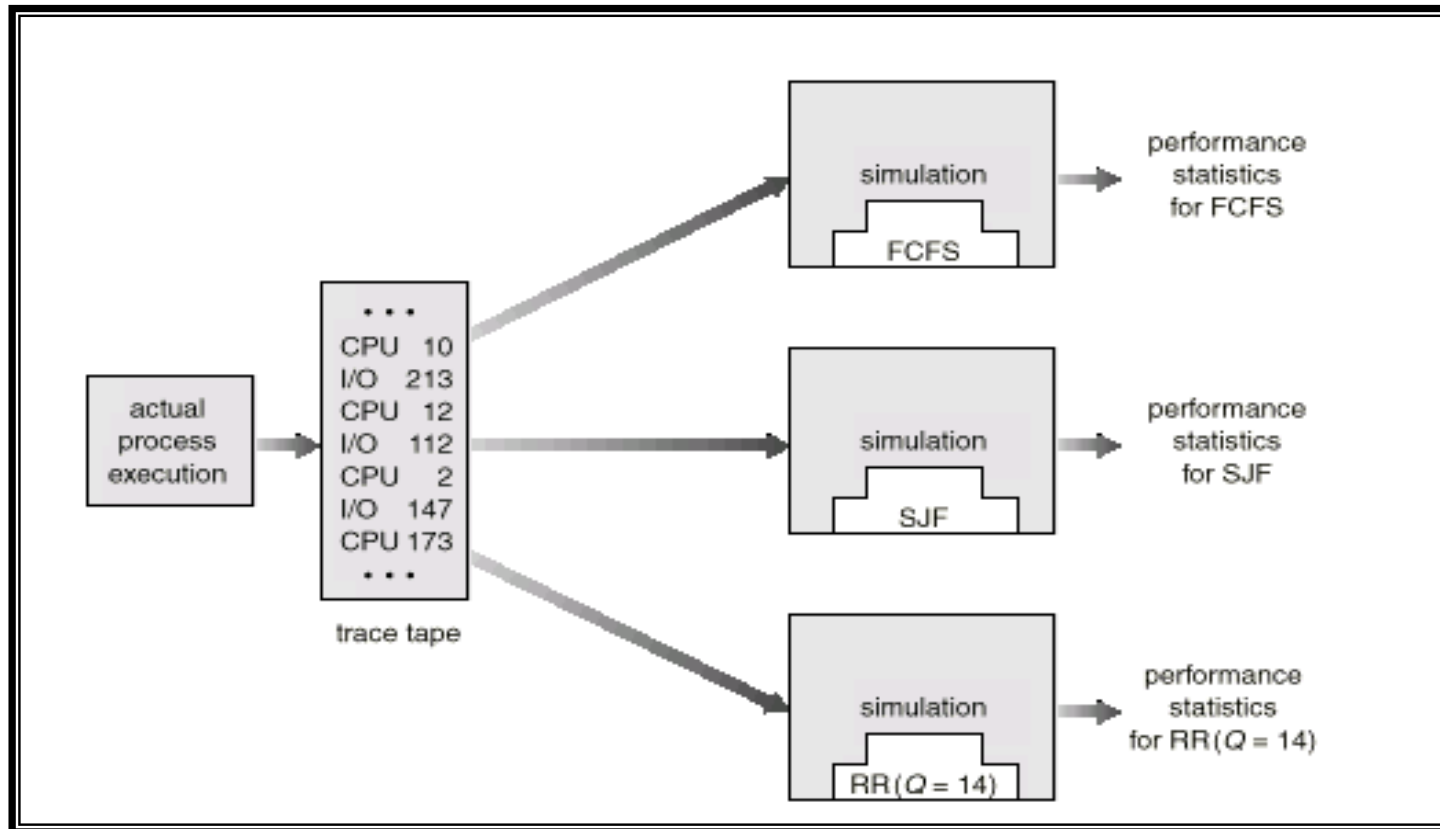
## ■ Implementation

- ✓ Direct implementation of the system under test, with appropriate benchmarks

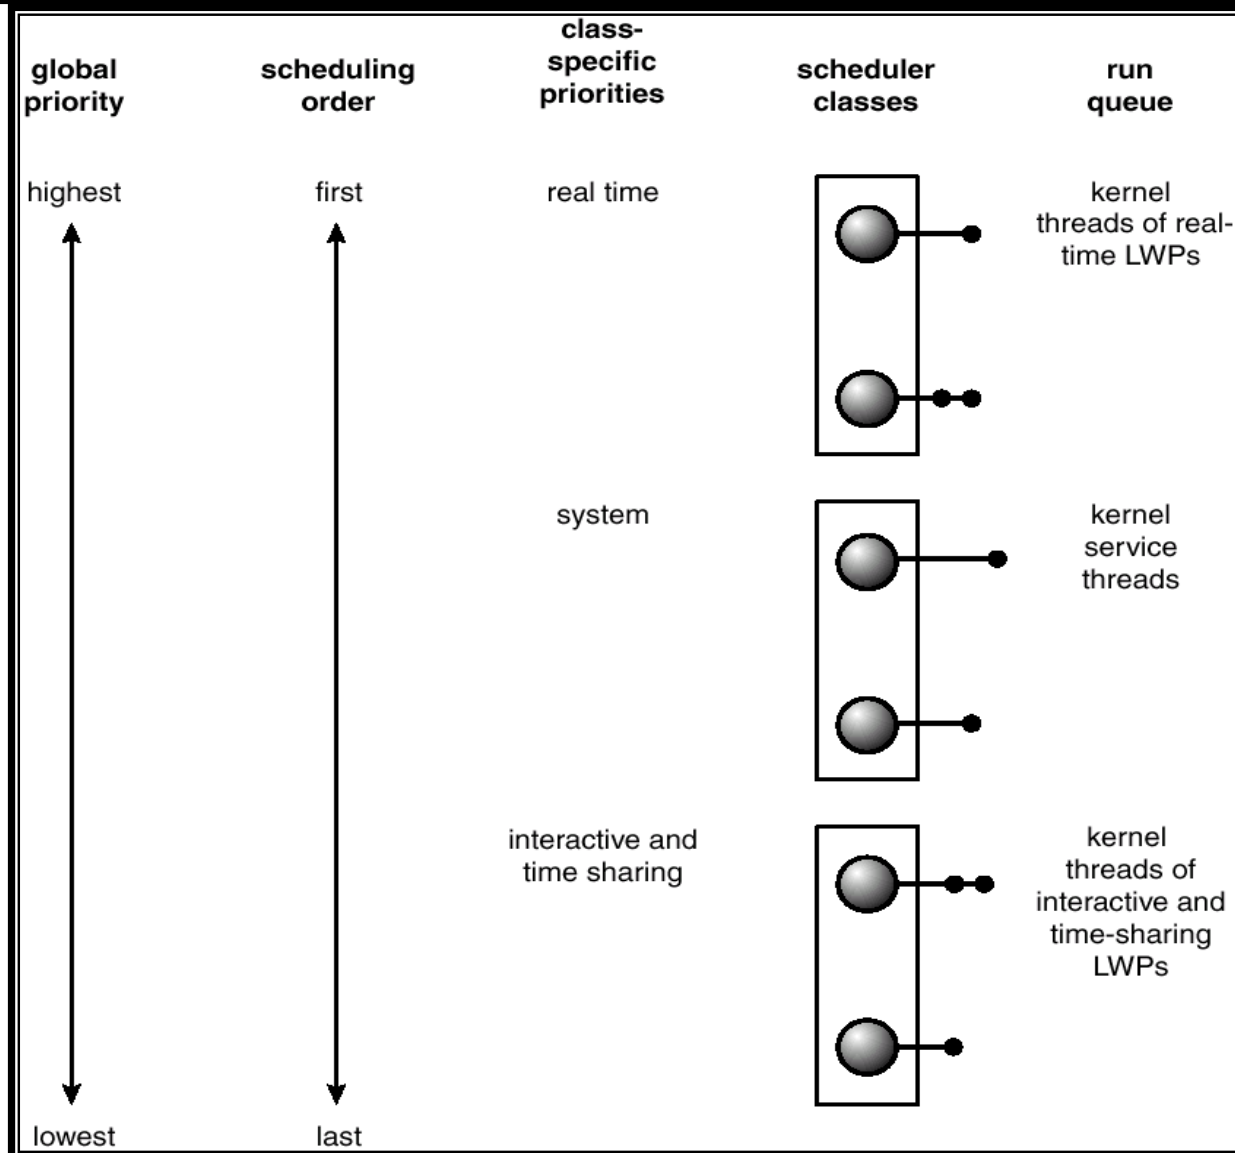




# Evaluation of CPU Schedulers by Simulation



# Solaris 2 Scheduling



# Windows XP Priorities

---

	real-time	high	above normal	normal	below normal	idle priority
time-critical	31	15	15	15	15	15
highest	26	15	12	10	8	6
above normal	25	14	11	9	7	5
normal	24	13	10	8	6	4
below normal	23	12	9	7	5	3
lowest	22	11	8	6	4	2
idle	16	1	1	1	1	1



## ■ Scheduling concept

- ✓ Selects a process from processes that are ready to execute and
- ✓ Allocates CPU to it (dispatcher)
- ✓ Preemptive vs. non-preemptive scheduling
  - Depends on whether the scheduler can interrupt/preempt a process and force a context switch

## ■ Scheduling goals

- ✓ Fairness
- ✓ Balance
- ✓ CPU utilization
- ✓ Throughput
- ✓ Turnaround time
- ✓ Waiting time
- ✓ Response time
- ✓ Meeting deadlines
- ✓ Avoiding starvation

## ■ Scheduling algorithms

- ✓ FCFS (First-Come First Served) or FIFO
- ✓ SJF (Shortest Job First)
- ✓ SRTF (Shortest Remaining Time First) or Preemptive SJF
- ✓ Priority scheduling
- ✓ Round Robin (RR)
- ✓ Multilevel Queue
- ✓ Multilevel Feedback Queue : canonical UNIX scheduler
- ✓ Real-time scheduling algorithms
  - Static priority: Rate-monotonic algorithm
  - Dynamic priority: EDF (Earliest Deadline First) algorithm

